

HF-Eval-S1/S2 Literature Map

Research reading artifact | fictional/internal | 2026-07-13

Analytical frame. This map describes only a fictional Helix Forge Lab dataset. It links the 2026-07-13 run R184 result for Model Alpha M-14 (F1 0.873) to the benchmark documentation needed to interpret it, rather than presenting a public dataset claim.

Dataset lineage / coverage

HF-Eval-S1/S2 v3.2 is documented as a deliberately synthetic two-part benchmark: S1 has 1,600 records and S2 has 800, for 2,400 total examples.

Label and slice questions

The five labels retrieve, route, refine, resolve, and abstain should be examined both overall and in ambiguity-focused slices.

Documentation gaps

The fictional card still needs explicit construction rules for prompt ambiguity, label adjudication, and the intended role of S2 as a stress partition.

Working matrix

Map node	Known coverage	Open documentation question
S1	1,600 synthetic records	How were ambiguity levels balanced?
S2	800 synthetic records	Which stress pattern was emphasized?
Labels	Five named intents	Which pairs are near neighbors?

Dataset lineage flow

Synthetic prompt construction -> five-label adjudication -> S1 reference partition (1,600) -> S2 stress partition (800) -> R184 evaluation.

Coverage gaps and documentation register

Gap	Documentation action	Owner role
Ambiguity construction	Describe prompt variants and intended confusion pairs	Dataset curator
Label adjudication	State fictional tie-break rules	Evaluation reviewer
S2 stress design	Name the targeted composition change	Benchmark steward

Map sources

HF-Eval Lineage Sketch, *Five-Intent Coverage Audit*, and *Synthetic Documentation Gaps Log* [synthetic/internal] define this action register.