

Model Alpha M-14: Prior Work Map

Research reading artifact | fictional/internal | 2026-07-13

Analytical frame. The map places Model Alpha M-14 within a fictional Helix Forge Lab research program. On 2026-07-13, run R184 recorded macro F1 0.873, precision 0.881, and recall 0.865; Model Beta B-07 recorded 0.846 macro F1. These figures anchor a dependency discussion, not a claim about external deployments.

Evidence families

Representation studies separate intent signals into semantic cues, routing context, and abstention indicators. The M-14 reading plan treats these as independently reviewable families.

Evidence / limitation matrix

Feature-removal records are evidence about local contribution; they do not establish a general account of language understanding.

Dependency map

Dataset documentation supports metric interpretation; metric interpretation determines which error slices deserve a second pass.

Working matrix

Family	Question	Constraint
Representation	Which cues distinguish route from resolve?	Keep feature groups removable.
Measurement	Which label masks weak recall?	Use macro F1 beside class readings.
Governance	Who can reproduce the comparison?	Use synthetic/internal records only.

Dependency and research-family map

Representation cues -> label evidence -> macro-F1 reading -> error-slice review. Dataset documentation supplies the assumptions at every arrow.

Unresolved evidence-gap queue

Gap	Why it matters	Next artifact
Route/resolve cue overlap	Could make the representation family look stronger than it is	Paired-case ledger
Abstain evidence boundary	Alters the meaning of precision improvement	Decision trace sample
S2 construction rationale	Determines whether stress behavior is interpretable	Dataset addendum

Citation trace

M-14 Feature Family Ledger [synthetic/internal] supplies the family names; *R184 Review Trace Map* [synthetic/internal] supplies the dependency arrows; *HF-Eval-S1/S2 v3.2 Design Card* [synthetic/internal] supplies the benchmark boundary.