Board A

1) Confu Grid (val split)

PRED →

TRUE ↓

	ALPH	BETO	GAMM	DELT
ALPH	72	9	6	3
BETO	12	63	10	5
GAMM	7	11	68	4
DELT	4	6	8	75

○ diag = good

□ off-diag = confusions

2) Calib Curve (reliability)

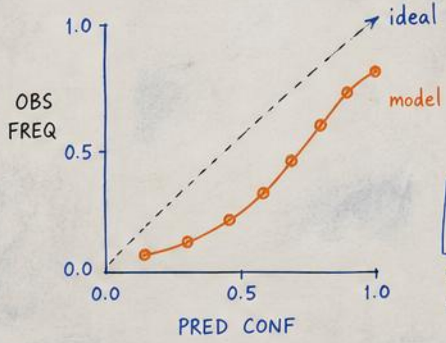

ECE-ish

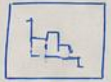

smaller is better

3) Latency vs Quality

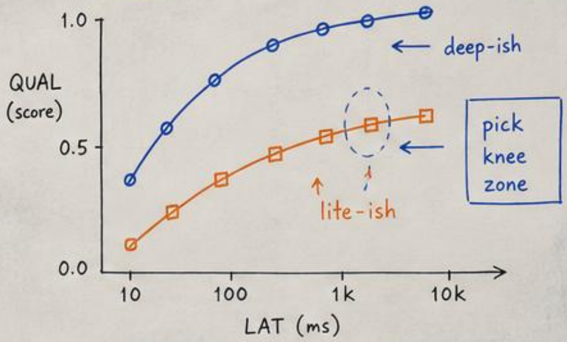

4) Decision Rules

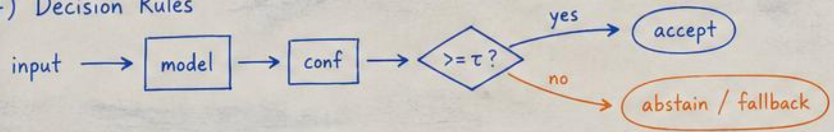

5) Threshold Sketch

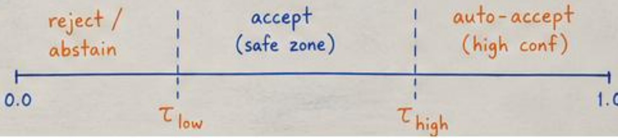

Board B

A) Pipeline Sketch

data mix → prep → model → score → decide → accept / abstain (review) / reject

model ↓ aux heads

B) Operating Regimes

more compute / more latency →

fast-lite	balanced	deep-ish
- lower qual	- mid qual	- highest qual
- cheap	- mid cost	- pricey
- broad use	- default?	- narrow use

C) Conf vs Threshold

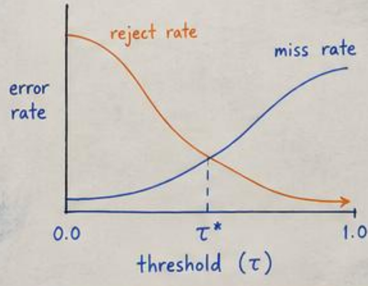

τ^* = sweet spot?

Goal: good qual where it's safe, abstain when unsure, learn & improve.

D) Cost Accounting (toy)

$$E[\text{cost}] = C_{run} + P_{abstain} \cdot C_{rev} + P_{miss} \cdot C_{miss}$$

minimize E[cost]

E) Notes / Ideas

- try temp-scaling or beta-ish calibration
- per-class τ? ⇒ adaptive
- collect hard cases for next round